



INDIGO - DataCloud

# PaaS-level Components

Alfonso Pérez, Miguel Caballer, Germán Moltó, Marica Antonacci, Lukasz Dutka

Universitat Politècnica de València, INFN-Bari,  
Cyfronet AGH



INDIGO-DataCloud is co-founded by the  
Horizon 2020 Framework Programme

# Index



- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData



# Index

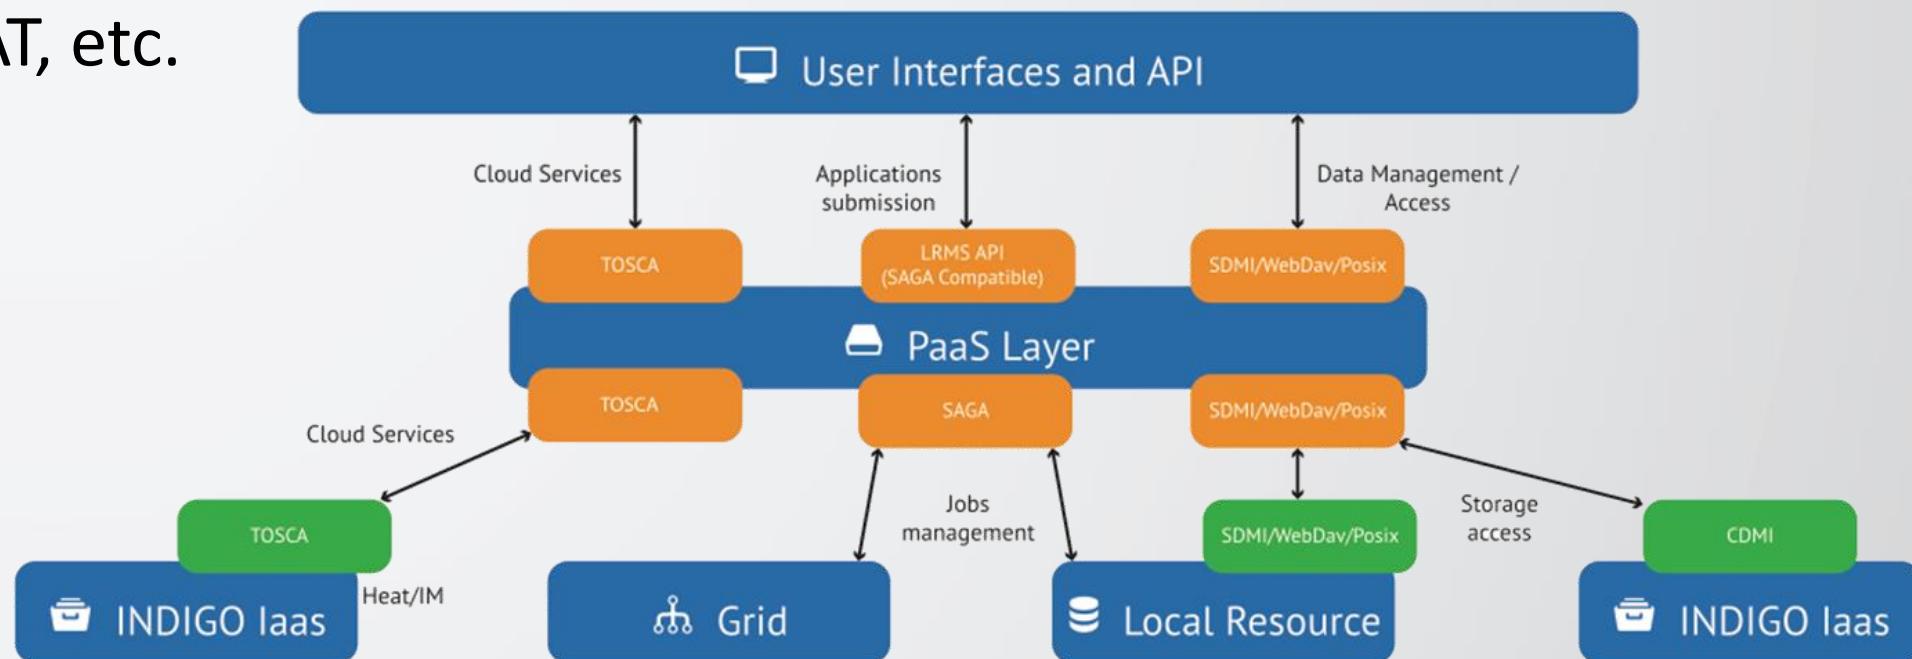


- **Indigo PaaS Overview**
- TOSCA
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData

# INDIGO advanced PaaS Platform



- INDIGO implements an advanced PaaS layer allowing scientific communities to exploit, in a powerful and high-level way, several heterogeneous computing and data e-infrastructure such as: IaaS Cloud, Helix Nebula, EGI Grid, EGI Federated Cloud, PRACE, HPC, EUDAT, etc.



# INDIGO PaaS: Key features

---



- Improved capabilities in the geographical exploitation of Cloud resources.
- Standard interface to access PaaS services.
  - INDIGO will use the TOSCA standard
- Support for data requirements in Cloud resource allocations.
  - Resources can be allocated where data is stored.
- Integrated use of resources coming from both public and private Cloud infrastructures

# INDIGO PaaS: Key features

---



- Distributed data federations supporting legacy applications as well as high level capabilities for distributed QoS and Data Lifecycle Management.
  - This includes for example remote Posix access to data.
- Transparent client-side import/export of distributed Cloud data.
  - This supports dropbox-like mechanisms for importing and exporting data from/to the Cloud. That data can then be easily ingested by Cloud applications through the INDIGO unified data tools.
- Support for distributed data caching mechanisms and integration with existing storage infrastructures.

# INDIGO PaaS: Key features

---



- Deployment, monitoring and automatic scalability of existing applications.
- Integrated support for high-performance Big Data analytics.
- Support for dynamic and elastic clusters of resources.
  - batch systems on-demand (such as HTCondor or Torque)
  - extensible application platforms (such as Apache Mesos) capable of supporting both application execution and instantiation of long-running services.

# Index



- Indigo PaaS Overview
- **TOSCA**
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData



# Standards

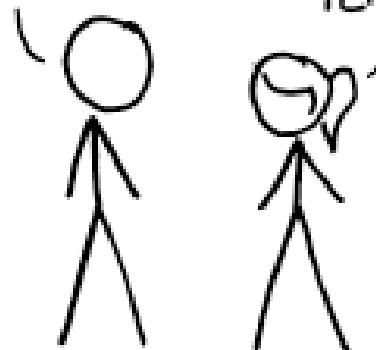


## HOW STANDARDS PROLIFERATE:

(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC)

SITUATION:  
THERE ARE  
14 COMPETING  
STANDARDS.

14?! RIDICULOUS!  
WE NEED TO DEVELOP  
ONE UNIVERSAL STANDARD  
THAT COVERS EVERYONE'S  
USE CASES.



YEAH!

SOON:

SITUATION:  
THERE ARE  
15 COMPETING  
STANDARDS.

# TOSCA

---

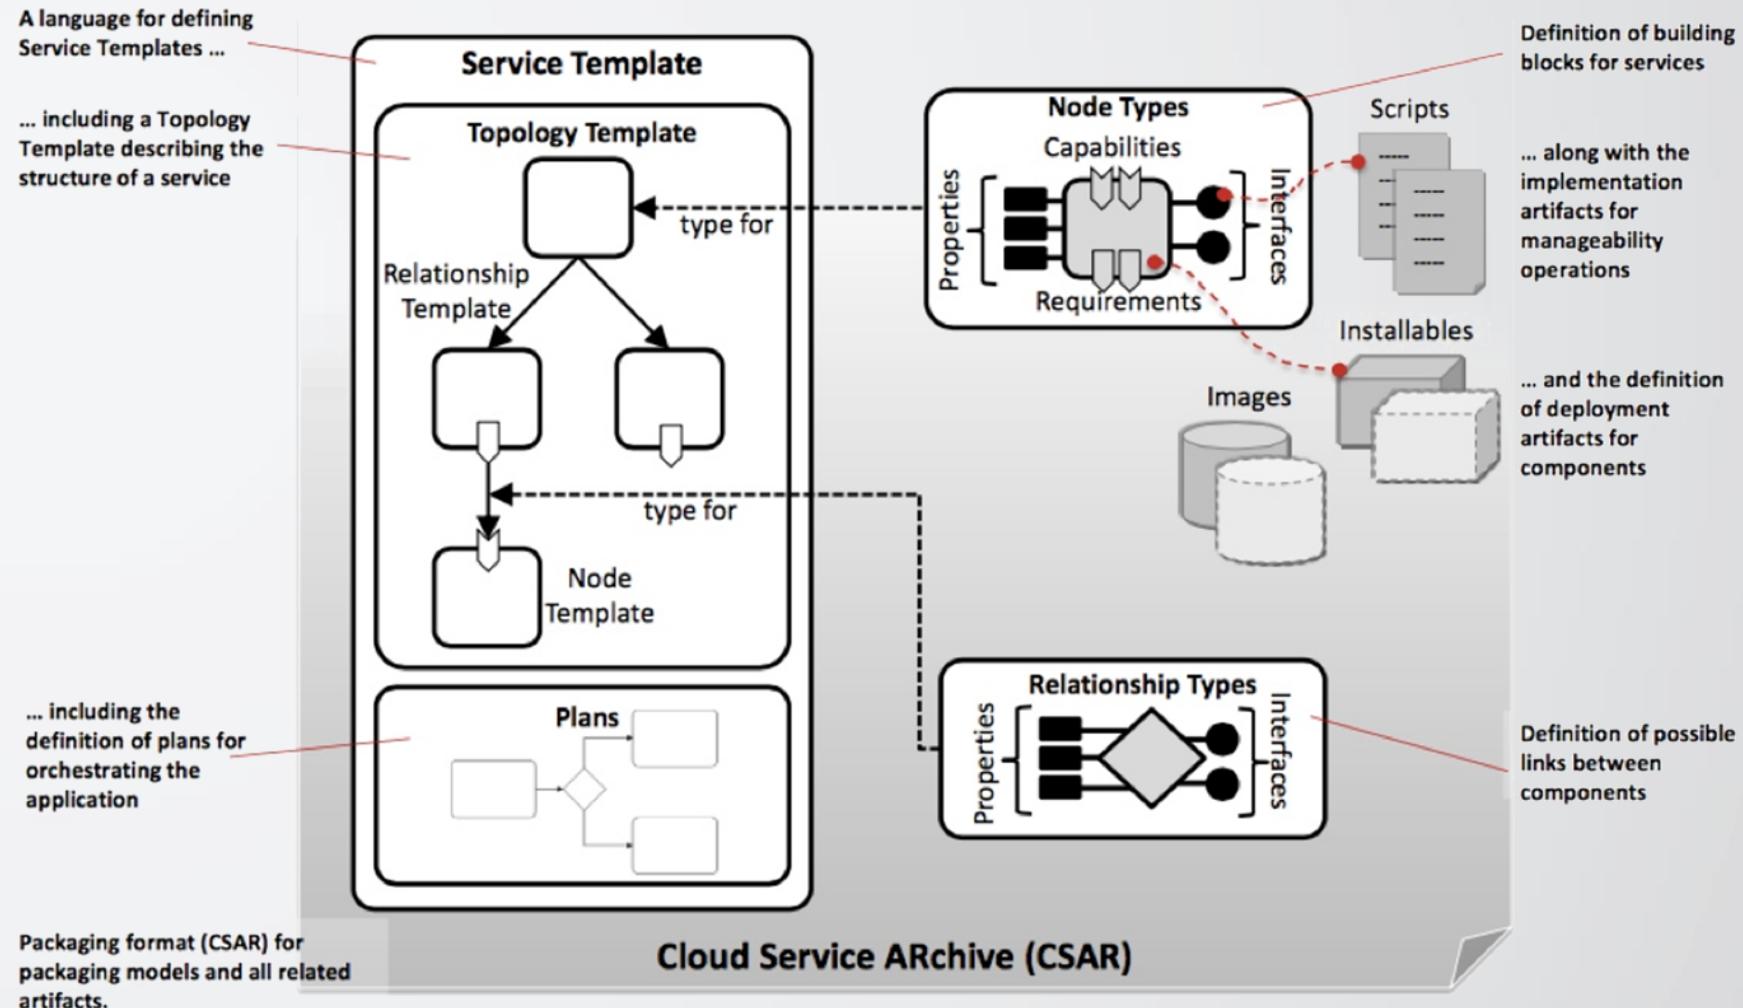


- Topology and Orchestration Specification for Cloud Applications
- Standardizes the language to describe
  - The structure of an ITService (its topology model)
  - How to orchestrate operational behavior (plans such as build, deploy, patch, shutdown, etc.)
    - Leveraging the BPMN standard
  - Declarative model that spans applications, virtual and physical infrastructure

# TOSCA in a nutshell



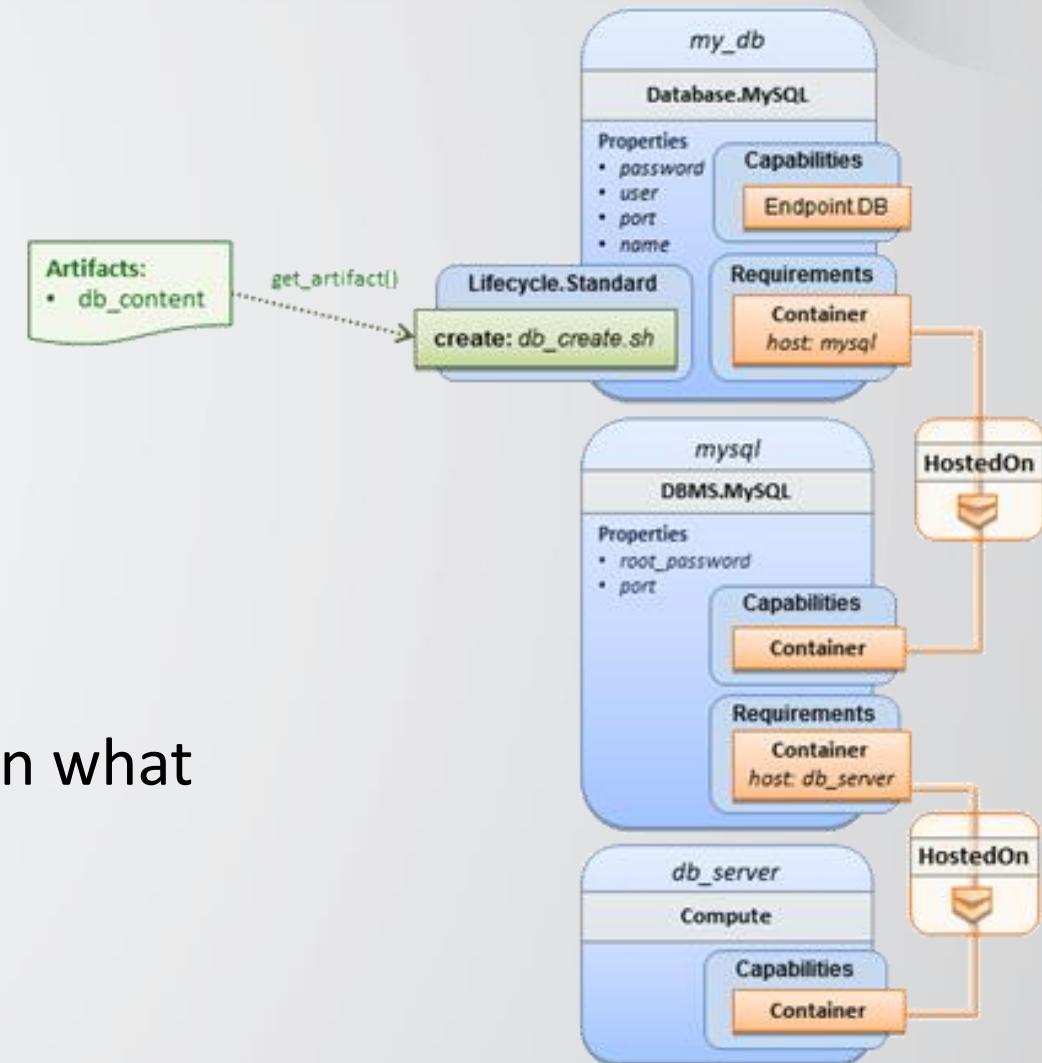
## OASIS Topology and Orchestration Specification for Cloud Applications



# TOSCA Topology



- 3 layers
  - Infrastructure (Cloud or DC objects)
  - Platform or Middleware (App containers)
  - Application modules, schemas and configurations
- Relationships between components:
  - What's hosted on what or installed on what
  - What's connected to what



# TOSCA Topology

---



- Components in the topology are called Nodes
- Each Node has a Type (e.g. Host, BD, Web server).
  - The Type is abstract and hence portable
  - The Type defines Properties and Interfaces
- An Interface is a set of hooks (named Operations)
- Nodes are connected to one another using Relationships
- Both Node Types and Relationship Types can be derived

# Node Type Examples



- Defines properties as YAML maps
- Might defines capabilities (What it can provide to other nodes)

```
tosca.nodes.DBMS
  derived_from: tosca.nodes.SoftwareComponent
  properties:
    dbms_root_password:
      type: string
      description: the root password for the DBMS service
    dbms_port:
      type: integer
      description: the port the DBMS service will listen to for data and requests
  capabilities:
    host:
      type: Container
      containee_types: [ tosca.nodes.Database ]
```

# Node Type Examples



- Might Define Requirements (what it needs from other nodes)

```
tosca.nodes.Database:  
  derived_from: tosca.nodes.Root  
  properties:  
    db_user:  
      type: string  
      description: user account name for DB administration  
    db_password:  
      type: string  
      description: the password for the DB user account  
    db_port:  
      type: integer  
      description: the port the underlying database service will listen to data  
    db_name:  
      type: string  
      description: the logical name of the database  
  requirements:  
    - host: tosca.nodes.DBMS  
  capabilities:  
    - database_endpoint: tosca.capabilities.DatabaseEndpoint
```

# Node Template



- An instance of a type (like Object to Class)
- Has specific properties
- Has artifacts:
  - What to install
  - How to install (mapped to interface hooks)
- Has requirements and capabilities (or relationships)

# Topology Template Example



```
tosca_definitions_version: tosca_simple_yaml_1_0
description: Template for deploying a single server with predefined properties.

topology_template:
  inputs:
    cpus:
      type: integer
      description: Number of CPUs for the server.
      constraints:
        - valid_values: [ 1, 2, 4, 8 ]
  node_templates:
    my_server:
      type: tosca.nodes.Compute
      capabilities:
        # Host container properties
        host:
          properties:
            # Compute properties
            num_cpus: { get_input: cpus }
            mem_size: 2048 MB
            disk_size: 10 GB
  outputs:
    server_ip:
      description: The private IP address of the provisioned server.
      value: { get_attribute: [ my_server, private_address ] }
```

# INDIGO custom types



- Extend the TOSCA normative types
  - E.g.: [tosca.nodes.indigo.GalaxyPortal](#), [tosca.nodes.indigo.GalaxyWN](#)

```
tosca.nodes.indigo.GalaxyPortal:  
  derived_from: tosca.nodes.WebServer  
  properties:  
    admin:  
      type: string  
      description: email of the admin user  
      default: admin@admin.com  
      required: false  
    admin_api_key:  
      type: string  
      description: key to access the API with admin role  
      default: not_very_secret_api_key  
      required: false  
    user:  
      type: string  
      description: username to launch the galaxy daemon  
      default: galaxy  
      required: false  
    install_path:  
      type: string  
      description: path to install the galaxy tool  
      default: /home/galaxy/galaxy  
      required: false  
  requirements:  
    - lrms:  
        capability: tosca.capabilities.indigo.LRMS  
        node: tosca.nodes.indigo.LRMS.FrontEnd  
        relationship: tosca.relationships.HostedOn  
  artifacts:  
    galaxy_role:  
      file: indigo-dc.galaxy  
      type: tosca.artifacts.AnsibleGalaxy.role  
  interfaces:  
    Standard:  
      configure:  
        implementation: https://raw.githubusercontent.com/indigo-dc/tosca-types/master/artifacts/galaxy/galaxy_install.yml  
  inputs:  
    galaxy_install_path: { get_property: [ SELF, install_path ] }  
    galaxy_user: { get_property: [ SELF, user ] }  
    galaxy_admin: { get_property: [ SELF, admin ] }  
    galaxy_admin_api_key: { get_property: [ SELF, admin_api_key ] }  
    galaxy_lrms: { get_property: [ SELF, lrms, type ] }
```

Ansible [playbook](#)

# A Sample tosca template: kepler



```
tosca_definitions_version: tosca_simple_yaml_1_0  
  
imports:  
- indigo_custom_types: https://raw.githubusercontent.com/indigo-dc/tosca-  
types/master/custom_types.yaml  
  
description: TOSCA template for deploying an instance for Kepler  
  
topology_template:  
  
inputs:  
  
    number_cpus:  
        type: integer  
  
        description: number of cpus required for the instance  
  
        default: 1  
  
    memory_size:  
        type: string  
  
        description: ram memory required for the instance  
  
        default: 1 GB  
  
node_templates:  
  
    kepler:  
        type: tosca.nodes.indigo.Kepler  
  
        requirements:  
        - host: kepler_server  
  
Computing requirements
```

```
kepler_server:  
    type: tosca.nodes.indigo.Compute  
  
    capabilities:  
  
        endpoint:  
            properties:  
                network_name: PUBLIC  
                ports:  
                    vnc_port:  
                        protocol: tcp  
                        source: 5900  
  
        host:  
            properties:  
                num_cpus: { get_input: number_cpus }  
                mem_size: { get_input: memory_size }  
  
        os:  
            properties:  
                type: linux  
                distribution: ubuntu  
                version: 14.04  
                image: one://onecloud.i3m.upv.es:2633/67  
  
    Network requirements  
  
outputs:  
  
    instance_ip:  
        value: { get_attribute: [ kepler_server, public_address, 0 ] }  
  
    instance_creds:  
        value: { get_attribute: [ kepler_server, endpoint.credential, 0 ] }  
  
OS requirements  
  
Image registered in the Cloud site
```

# Index



- Indigo PaaS Overview
- TOSCA
- **IM**
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData



# Landing page



IM: Infrastructure Manager

Features   Learn More   Documentation   Use cases   About

*Automated DevOps for your Virtual Infrastructures*

**Deploy Customized Virtual  
Infrastructures on Multi-Clouds**

DOWNLOAD

TRY IT NOW!



<http://www.grycap.upv.es/im>

# Features



IM: Infrastructure Manager

Features   Learn More   Documentation   Use cases   About

## IM MAIN FEATURES



### Multi-Backend

Deploy on on-premises, public and scientific Clouds, and container orchestration platforms.



### Extensible plugins

Plugins available for [OpenNebula](#), [OCCI](#), [Amazon EC2](#), [Google Cloud Platform](#), [Microsoft Azure](#), [Docker](#), [Kubernetes](#), [FogBow](#), [OpenStack](#), [libvirt](#) and [EGI Federated Cloud](#).



### Hybrid Infrastructures

Deploy virtual infrastructures that span across multiple providers.



### Embrace DevOps

Powered by [Ansible](#), the IM provides recipes for common deployments (Hadoop clusters, etc.).



### Interfaces

Featuring a CLI, a web GUI, an XML-RPC service API and a REST API.



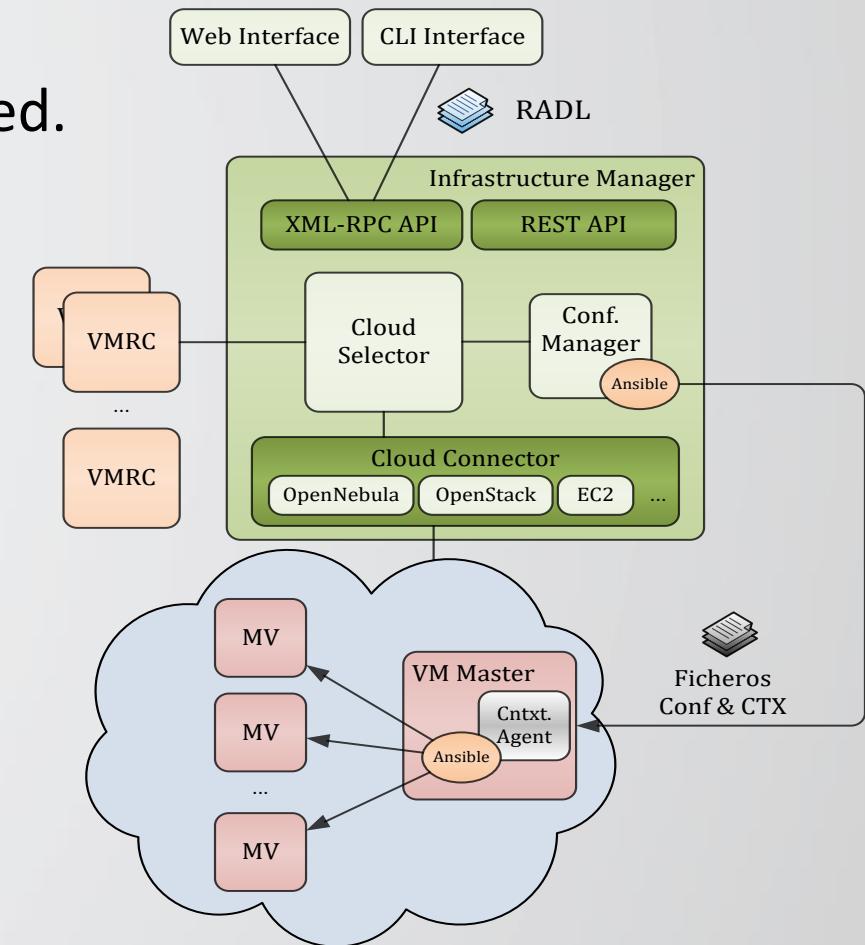
### Open Source

Developed in [Python](#), hosted in [GitHub](#) and distributed under the [GNU GPL v3.0](#). Docker images for the [IM server](#) and [IM web](#) are available.

# Introduction



- General platform to deploy on demand customizable virtual computing infrastructures.
  - With the precise software configuration required.
  - Allow to deploy any kind of complex infrastructure.
  - Share Infrastructure descriptions.
  - No need of pre-baked VMIs.
  - The same complex infrastructure can be deployed both on-premises and in a public Cloud.



# Cloud providers



- It supports a wide range of cloud providers and other computing back-ends :

- **Public:** Amazon Web Services (AWS),

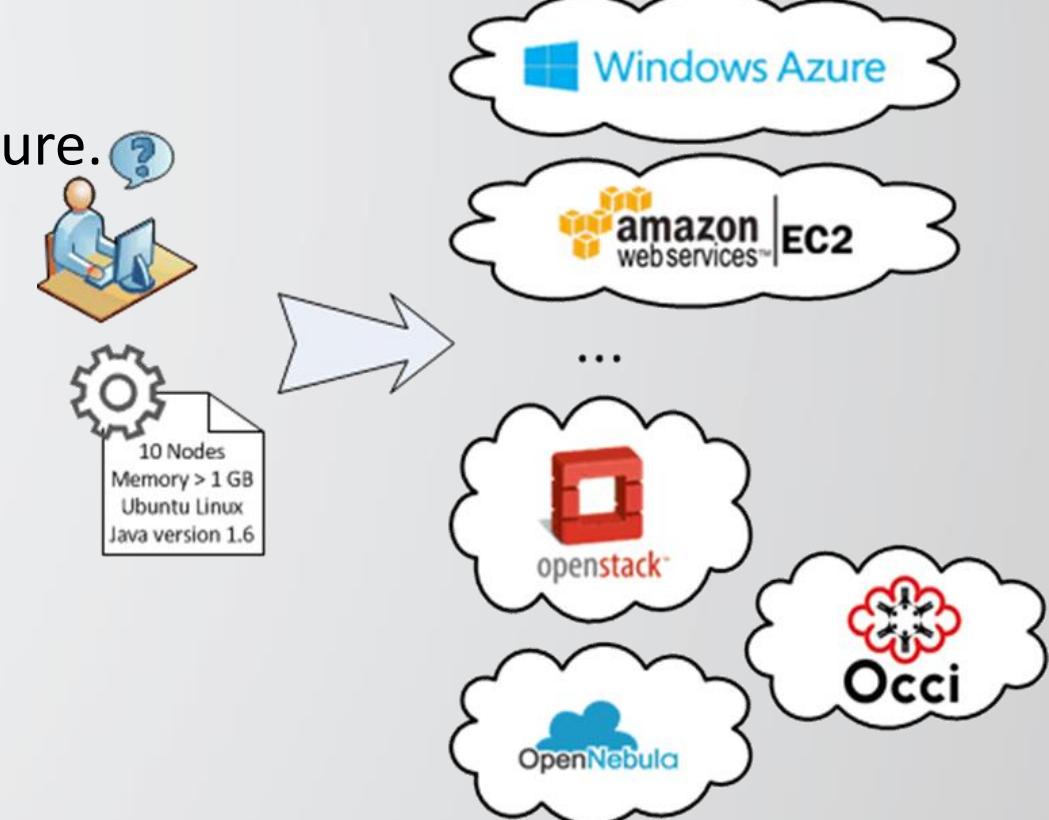
- Google Cloud Platform (GCP), Microsoft Azure.

- **On-premises:** OpenNebula, OpenStack, libvirt.

- **Federated:** EGI FedCloud (OCCI), FogBow.

- **Containers:** Docker, Kubernetes

- The list above can be easily extended by plugins.



# Main features

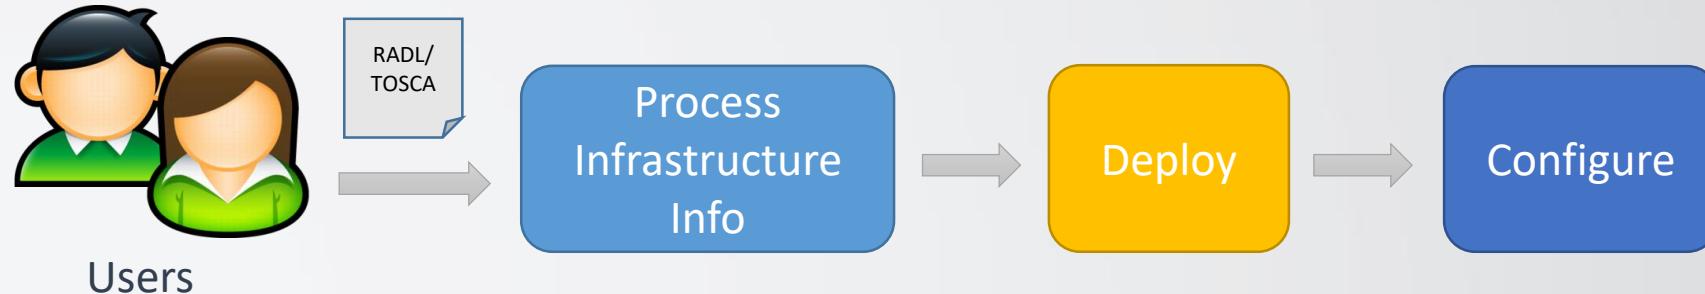


- It features **DevOps** capabilities.
  - Based on Ansible.
  - Provides recipes for common deployments.
  - Also supporting cloud-init scripts.
- IM works as a service that offers several **interfaces**:
  - XML-RPC and REST APIs.
  - Command-line application.
  - Web-based GUI.
- It is distributed under a **GNU GPL v3.0** open source license and its source code is available on **GitHub**.



<https://github.com/grycap/im>

# Working Scheme (I)



- The user can provide an RADL or TOSCA documents as input to the IM, describing the infrastructure:
  - **RADL:**
    - Resource and Application Description Language.
    - High level Language to define virtual infrastructures and Specify VM requirements.
  - **TOSCA:**
    - OASIS Standard
    - Open standard language to model application architectures to be deployed on a Cloud.

# RADL Document



```
ansible <ansible_host_id> (<features>  
  
network <network_id> (<features>  
  
system <system_id> (<features>  
  
configure <configure_id> (<Ansible recipes>  
  
contextualize [max_time] ( system <system_id>  
  
configure <configure_id> [step <num>] ... )  
  
deploy <system_id> <num> [<cloud_id>]
```

The keywords `ansible`, `network`, `system` and `configure` assign some *features* or *recipes* to an identity `<id>`. The features are a list of constraints separated by `and`, and a constraint is formed by `<feature name> <operator> <value>`.

```
network net (inbound = 'no')  
  
system node (  
    cpu.arch = 'x86_64' and  
    cpu.count = 1 and  
    memory.size >= 512M and  
    net_interface.0.connection = 'net' and  
    disk.0.os.name = 'linux'  
)  
  
configure node (@begin  
---  
- tasks:  
  - user: name=user1 password=1234  
@end  
)  
  
contextualize (  
    system node configure node  
)  
  
deploy node 1
```

# Using IM: A simple node with Ansible



```
network publica (outbound = 'yes')

system front (
    instance_type = 'extra-large' and
    net_interface.0.connection = 'publica' and
    disk.0.os.name = 'linux' and
    disk.0.image.url = 'https://one.upv.es:9856/495' and
    disk.0.os.credentials.username = 'cloudadm' and
    disk.0.os.credentials.password = 'xxx'
)
configure front (
@begin
    - name: Install some packages
        action: yum pkg(pkg1,pkg2) state=installed
@end
)
```

# A simple node with cloud-init



```
network publica (outbound = 'yes')

system front (
    instance_type = 'extra-large' and
    net_interface.0.connection = 'publica' and
    disk.0.os.name = 'linux' and
    disk.0.image.url =
        'https://fc-one.i3m.upv.es:11443/uuid\_image\_for\_egi\_centos\_6\_centos6kvm\_im154\_170' and
    disk.0.os.credentials.username = 'cloudadm' and
    disk.0.os.credentials.password = 'xxx'
)

configure front (
@begin
    #!/bin/bash
    yum install -y pkg1 pkg2
@end
)
contextualize (
    system front configure front with cloud_init
)
```

# A Sample tosca template: kepler



```
tosca_definitions_version: tosca_simple_yaml_1_0  
  
imports:  
- indigo_custom_types: https://raw.githubusercontent.com/indigo-dc/tosca-  
types/master/custom_types.yaml  
  
description: TOSCA template for deploying an instance for Kepler  
  
topology_template:  
  
inputs:  
  
    number_cpus:  
        type: integer  
  
        description: number of cpus required for the instance  
  
        default: 1  
  
    memory_size:  
        type: string  
  
        description: ram memory required for the instance  
  
        default: 1 GB  
  
node_templates:  
  
    kepler:  
        type: tosca.nodes.indigo.Kepler  
  
        requirements:  
        - host: kepler_server  
  
Computing requirements
```

```
kepler_server:  
    type: tosca.nodes.indigo.Compute  
  
    capabilities:  
  
        endpoint:  
            properties:  
                network_name: PUBLIC  
                ports:  
                    vnc_port:  
                        protocol: tcp  
                        source: 5900  
  
        host:  
            properties:  
                num_cpus: { get_input: number_cpus }  
                mem_size: { get_input: memory_size }  
  
        os:  
            properties:  
                type: linux  
                distribution: ubuntu  
                version: 14.04  
                image: one://onecloud.i3m.upv.es:2633/67  
  
    Network requirements  
  
outputs:  
  
    instance_ip:  
        value: { get_attribute: [ kepler_server, public_address, 0 ] }  
  
    instance_creds:  
        value: { get_attribute: [ kepler_server, endpoint.credential, 0 ] }  
  
OS requirements  
  
Image registered in the Cloud site
```

# Working Scheme (II)

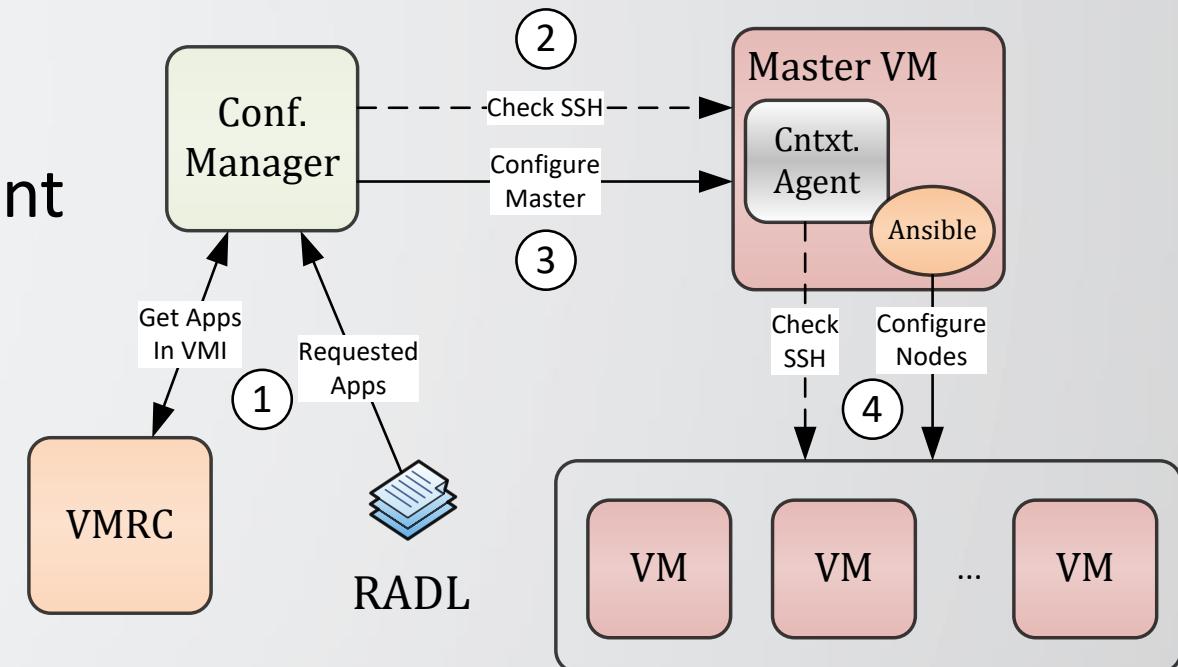


- The user specifies the image (or list of images) to use.
  - URI naming convention to abstract from cloud provider:
    - one://server:port/image-id
    - ost://server:port/ami-id
    - aws://region/ami-id
    - <site end-point>/<image-id>
  - The IM can also contact the VMRC (<http://www.grycap.upv.es/vmrc>) to get a list with the most suitable VMIs
    - In INDIGO-DataCloud, the image information is obtained from the CMDB.
- Then, the IM obtains the list of IaaS providers available to the user.
  - From the credentials provided by the user.
- Finally, it contacts the IaaS provider selected and deploys the infrastructure.

# Contextualization process: with ansible



1. SSH connection to the Master VM
  - A GNU/Linux-based VM with a public IP
2. Configure Master VM
  - Install and configure Ansible
3. Launch Contextualization Agent
  - Check SSH from VMs
  - Call Ansible



# Client Tools



- Command Line Tool (CLI).

```
usage: client.py [-u|--xmlrpc-url <url>] [-a|--auth_file <filename>] operation op_parameters
options:
  --version           show program's version number and exit
  -h, --help          show this help message and exit
  -a AUTH_FILE, --auth_file=AUTH_FILE
                      File with authentication data
  -u XML-RPC, --xmlrpc-url=XML-RPC
                      URL of the InfrastructureManager service.
```

## Operations:

list	
create	<radl_file>
destroy	<inf_id>
getinfo	<inf_id> [radl_attribute]
getradl	<inf_id>
getcontmsg	<inf_id>
getstate	<inf_id>
getvminfo	<inf_id> <vm_id> [radl_attribute]
getvmcontmsg	<inf_id> <vm_id>

addresource	<inf_id> <radl_file> [ctxt flag]
removeresource	<inf_id> <vm_id> [ctxt flag]
alter	<inf_id> <vm_id> <radl_file>
start	<inf_id>
stop	<inf_id>
reconfigure	<inf_id> [<radl_file>] [vm_list]
startvm	<inf_id> <vm_id>
stopvm	<inf_id> <vm_id>
sshvm	<inf_id> <vm_id>
getversion	

# Client-side Tools: Web



- Publicly-available web interface (also open-sourced):
  - Easily deploy infrastructures from a web browser
  - Share RADL/TOSCA documents.
  - <http://servproject.i3m.upv.es/im/>

The screenshot shows the Infrastructure Manager web interface. On the left, there's a sidebar with links for Infrastructures, Credentials (which is selected), RADLs, EC3, and Admin. The main content area has a title "Infrastructure Manager > Credentials" and a table with columns: ID, Type, Host, Edit, Delete, Enabled, and Order. It lists three entries: one (OpenNebula, host onecloud.i3m.upv.es:2633), ec2 (amazon|EC2, host blank), and VMRC (VMRC, host http://servproject.i3m.upv.es:8080/vmrc/vmrc). At the bottom, it says "Showing 1 to 4 of 4 entries". The footer includes the GRyCAP logo and university information.

The screenshot shows the Infrastructure Manager web interface. The sidebar has links for Infrastructures, Credentials, RADLs (selected), EC3, and Admin. The main content area has a title "Infrastructure Manager" and a table titled "List" showing entries for Infrastructures, Credentials, RADLs, EC3, and Admin. The RADLs section shows two entries: "hadoop" (Hadoop Cluster 5 WNs) and "hadoop-cluster" (Deploys a 3-node Hadoop Cluster). The footer includes the GRyCAP logo and university information.

# APIs to be consumed by clients



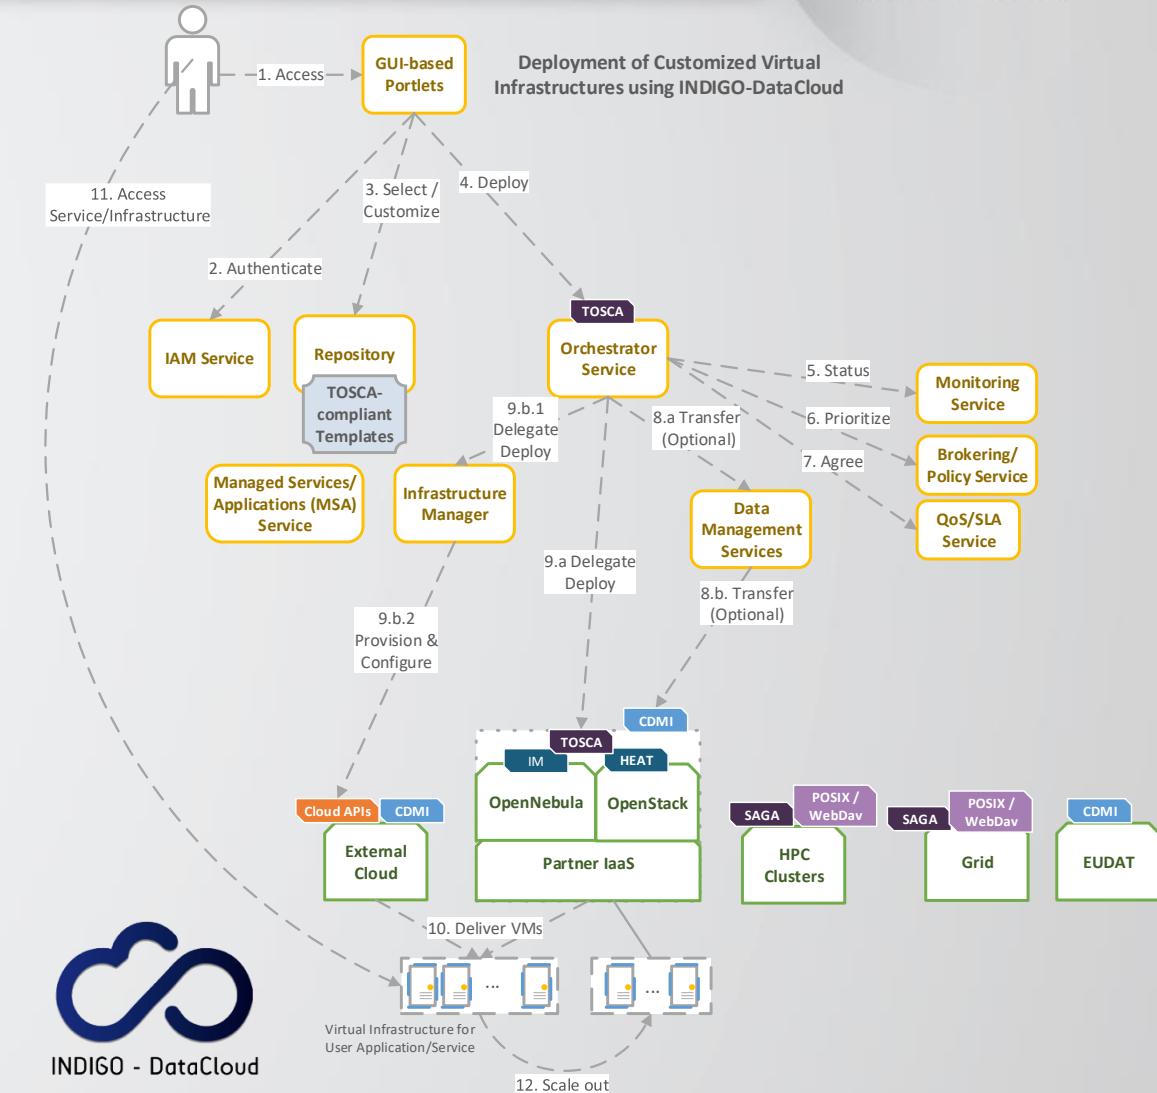
- **XML-RPC API**
  - API that follows the XML-RPC specification.
- **REST API**
  - IM Service can be accessed through a REST(ful) API
- More info: <http://www.grycap.upv.es/im/documentation.php>

HTTP method	/infrastructures	/infrastructures/<infId>	/infrastructures/<infId> /vms/<vmId>
GET	List the infrastructure IDs.	List the virtual machines in the infrastructure <code>infId</code>	Get information associated to the virtual machine <code>vmId</code> in <code>infId</code> .
POST	Create a new infrastructure based on the RADL posted	Create a new virtual machine based on the RADL posted.	
PUT			Modify the virtual machine based on the RADL posted.
DELETE		Undeploy all the virtual machines in the infrastructure.	Undeploy the virtual machine.

# Where is the im used? - INDIGO-DataCloud



- In the **INDIGO-DataCloud** project:
  - IM is a **key component** of the architecture:
    - Used at the PaaS Core to provide deployment of infrastructures to Cloud sites external to INDIGO-DataCloud (including public Cloud sites).
    - Used at IaaS to provide TOSCA-based deployment of infrastructures for OpenNebula sites.



# Where is the IM used? - EGI



- The IM is used in the **VMOps Dashboard of EGI**.
  - As the OCCI communication layer to create VM topologies.
  - TOSCA compatible.
  - [https://wiki.egi.eu/wiki/Federated\\_Cloud\\_AppDB\\_VMOps\\_Dashboard](https://wiki.egi.eu/wiki/Federated_Cloud_AppDB_VMOps_Dashboard)

A screenshot of the EGI Applications Database (AppDB) interface. The top navigation bar includes links for Home, Software Marketplace, Cloud Marketplace, and People. The main content area shows a listing for the "Infrastructure Manager (IM)". The listing includes a thumbnail image of a yellow and black cube, the title "Infrastructure Manager (IM)", a rating of "Excellent (Avg.: 5.00 by 1 vote)", and a brief description: "The IM is a tool that ease the access and the usability of IaaS clouds by automating the VMI selection, deployment, configuration, software installation, monitoring and update of Virtual Appliances." Below the listing are categories like Tools and Science Gateways, disciplines like Infrastructure Development, and tags like ::EC2 API, ::OpenNebula, ::OpenStack, and ::Spain. A sidebar on the right provides communication options: Join as a contact, Send a message, and Report a problem.

Home > Software > Science Gateways > Infrastructure Manager (IM)

Information Publications (3) Software Releases Comments & Ratings

Infrastructure Manager (IM) (id:880) [permalink]

Rate It: ★★★★★ Excellent (Avg.: 5.00 by 1 vote)

The IM is a tool that ease the access and the usability of IaaS clouds by automating the VMI selection, deployment, configuration, software installation, monitoring and update of Virtual Appliances.

Category: Tools • Science Gateways ▾

Disciplines: Infrastructure Development ▾

Tags: add ::EC2 API ::OpenNebula ::OpenStack ::Spain

<https://appdb.egi.eu/store/software/infrastructure.manager.im>



# Where is the IM used? - INDRA



- The **INDRA** company has integrated the IM in their GPaas Cloud platform.
  - It permits the coordinated deployment and configuration of middleware and applications within the context of a specific infrastructure.
  - <http://www.indracompany.com/en/node/64122>

A screenshot of the INDRA Cloud Computing website. At the top, there's a navigation bar with the INDRA logo, followed by links for INDUSTRIES, CONSULTING, TECHNOLOGY (underlined), OUTSOURCING (underlined), DIGITAL, and ABOUT INDRA. Below the navigation is a large heading "CLOUD COMPUTING". Underneath it, a blue banner features the text "PLATAFORMA IN CLOUD &gt; INFRASTRUCTURE MANAGER - GIM". At the bottom of the page, there's a paragraph describing the tool: "A tool for automating IT processes. It permits the coordinated deployment and configuration of middleware and applications within the context of a specific infrastructure." To the right of this text is a teal sidebar containing a profile picture of Alfonso Ríos Alonso and the text "Alfonso Ríos Alonso Cloud Computing". There are also icons for email and LinkedIn.

# Where is the im used? - EC3



- IM is a key component in **EC3**, a tool to deploy virtual hybrid elastic clusters that is integrated in the **EGI Access** service (for the long-tail of science).
  - EC3: <http://www.grycap.upv.es/ec3/>
  - EGI Access: <https://access.egi.eu/start>

The image displays two screenshots of the EC3 (Elastic Cloud Computing Cluster) website. The left screenshot shows the homepage with a dark blue background featuring a glowing blue light effect. It includes the text "Cluster as a Service" and "Deploy Virtual Elastic Clusters on the Cloud". Below this is a large "EC3" logo, and at the bottom are two yellow buttons: "DEPLOY YOUR CLUSTER!" and "LEARN MORE". The right screenshot shows a page titled "WHERE DO YOU WANT TO DEPLOY THE CLUSTER?". It asks users to provide valid credentials for the cloud provider and mentions the ability to deploy a hybrid cluster via the CLI. It lists four deployment options: Amazon Web Services (with a yellow icon), OpenNebula (with a blue icon), OpenStack (with a red icon), and EGI FedCloud (with a white icon featuring a green gear and blue dots). Each option has a brief description below it.

EC3: Elastic Cloud Computing Cluster

FEATURES LEARN MORE DEPLOY! CONTACT

*Cluster as a Service*

**Deploy Virtual Elastic Clusters on the Cloud**

DEPLOY YOUR CLUSTER! LEARN MORE

EC3: Elastic Cloud Computing Cluster

FEATURES LEARN MORE DEPLOY! CONTACT

**WHERE DO YOU WANT TO DEPLOY THE CLUSTER?**

You will need to provide valid credentials for the Cloud provider. Not sure if this is safe? Check the docs.

Wanted to deploy a hybrid cluster? You can do it with the CLI.

**Amazon Web Services™**  
Public Cloud provider

**OpenNebula**  
On-premises Cloud provider

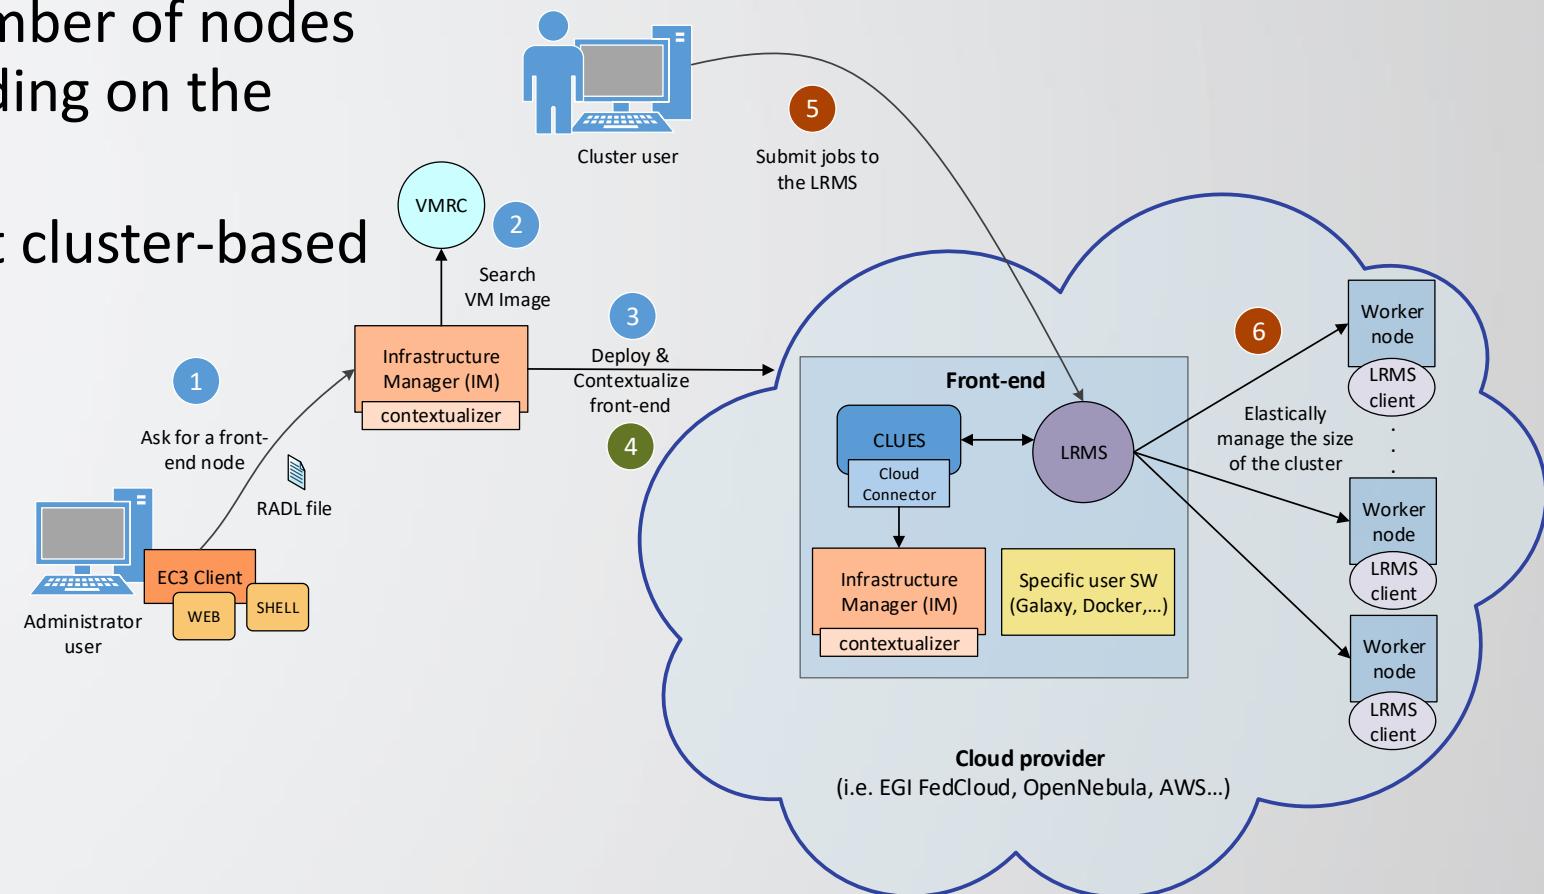
**openstack™ CLOUD SOFTWARE**  
On-premises Cloud provider

**EGI FedCloud**  
European Federated Cloud  
(See a case study here)

# Where is the im used? - EC3



- The IM provisions the front-end node of a virtual cluster where the number of nodes can grow and shrink (depending on the workload).
  - Cost and energy-efficient cluster-based computing.



# Index



- Indigo PaaS Overview
- TOSCA
- IM
- **CLUES**
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData



# Automatic Scaling Service



- Extends EC3 CLUES adding the interfaces required to interact with the INDIGO Orchestrator
  - Documentation: <http://www.grycap.upv.es/ec3>
- Implements the elasticity rules considering the state of the virtual cluster.
- The virtual cluster will deploy additional worker nodes as required, and integrate them on the LRMS without user intervention, in order to cope with increased workload of jobs. Worker nodes will be terminated when they are no longer required.
- Plugins are available for: SLURM, Torque/PBS, HTCondor, Mesos

# Index



- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- **Orchestrator**
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData



# Orchestrator Service

---

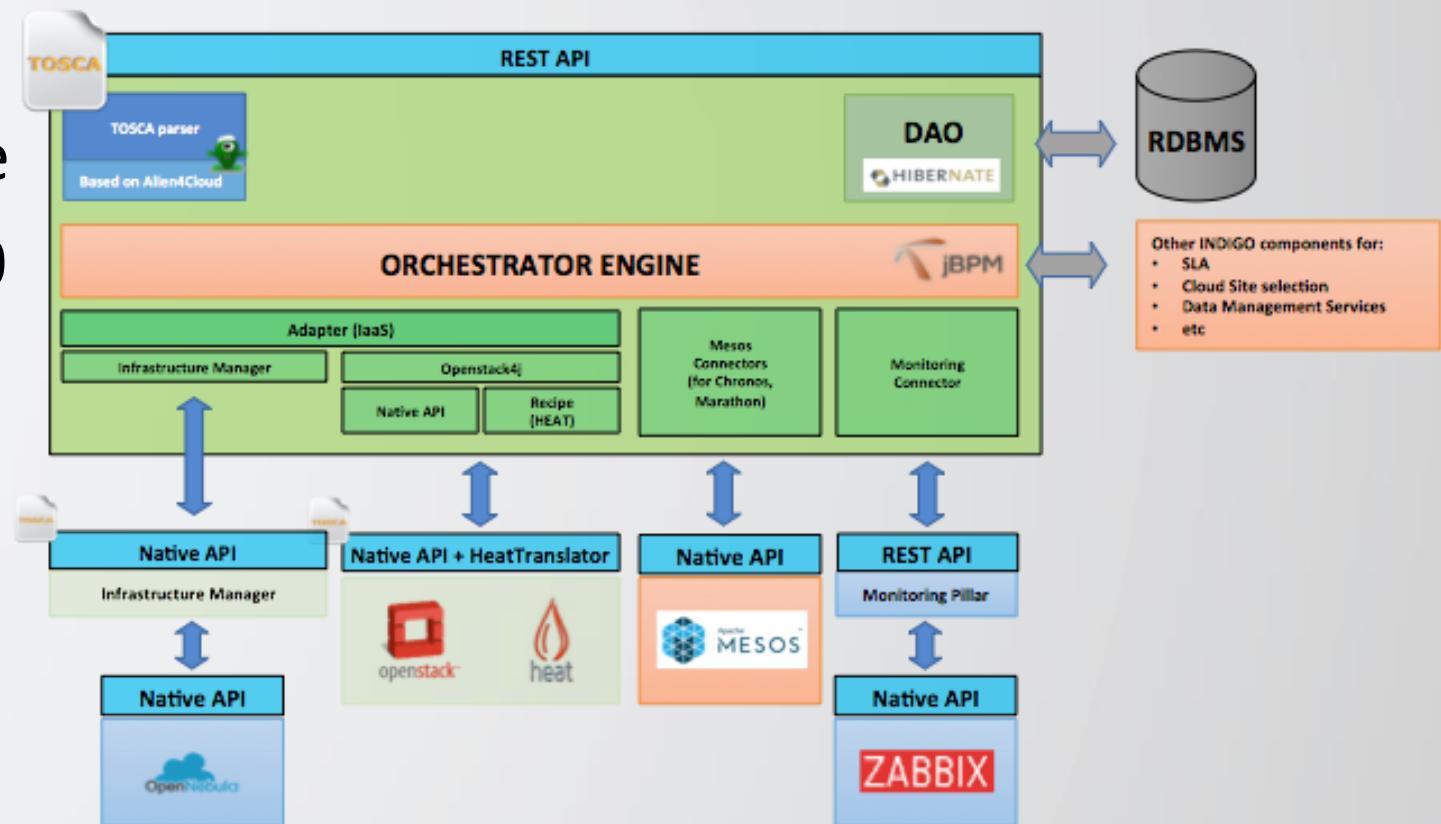


- The Orchestrator **coordinates the deployment process** over the IaaS platforms
- The Orchestrator collects all the information needed to deploy a service consuming others PaaS µServices APIs:
  - ***Monitoring Service***: get the capabilities of the underlying IaaS platforms and their resource availability;
  - ***QoS/SLA Service***: get the prioritized list of SLA per user/group
  - ***CloudProviderRanker*** (Rule Engine) Service: sort the list sites on the basis of rules defined per user/group/use-case;
  - ***Data Management Service***: get the status of the data files and storage resources needed by the service/application
- The orchestrator **delegates** the deployment to **IM**, **HEAT** or **Mesos** based on the TOSCA template and the list of sites.
- **Cross-site deployments** will also be possible.

# Orchestrator Service



- Built in **JBPM 6.1** (long-running workflow)
- Exposes **RESTful APIs**
- Supports the **TOSCA Simple Profile in YAML Version 1.0** specification



# Index



- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- Orchestrator
- **Cloud Provider Ranker**
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData



# Cloud Provider Ranker Service



- Provides information that will be consumed by the Orchestrator in order to properly coordinate the deployment of the required resource on the sites.
- A WEB Service providing REST APIs to rank Cloud Providers described by a JSON blob containing:
  - Total VirtualCPUs, total VirtualRAM
  - Total Virtual Ephemeral Disk (space for instances)
  - Total Virtual Disk (block storage, e.g. Cinder)
  - In use VCPU, in use VRAM, in use VDISK, in use VEphDisk
- Ranking algorithm implemented using the largely diffused **Drools Rule Engine** runtime framework

# Index



- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- **Monitoring Service**
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData

# Monitoring Service

---



- The Monitoring service provides a comprehensive REST API to gather information about
  - the **PaaS Core Services**
    - Using *Heapster* in the *Kubernetes* cluster.
  - the **customized virtual infrastructures**
    - Using *Zabbix agents* deployed inside VMs/containers
  - the **state of the sites** (leveraging EGI FedCloud approach)
- Monitored information has to be exposed via a REST API to be consumed by other services

# Index



- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- **SLAM**
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData

# QoS/SLA Service



- Allows the **handshake** between a **user** and a **site** on a given **SLA**
- Provides the Orchestrator/Ranker with the useful information for taking the decision on tasks scheduling according to the agreed and valid SLAs
- Describes the **QoS** that a specific user/group has both over a given site or generally in the PaaS as a whole; this includes a priority of a given users, the capability to access to different QoS at each site (Gold, Silver, Bronze services)

# Index



- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- **Accounting Service**
- CMDB
- Managed Services/Application (MSA) Deployment Service
- OneData



# Accounting Service

---



- Accounts for resource usage on the INDIGO PaaS and provides that data to other INDIGO-DataCloud services
  - QoS/SLA service will use information gathered by the Accounting service (and the monitoring pillar) to monitor SLA violation
  - Usage data is extracted from the system where the resources are used and sent to a central repository
  - The repository aggregates the data from across the infrastructure to produce totals based on a number of fields - such as user, site, month, year, etc.

# Index

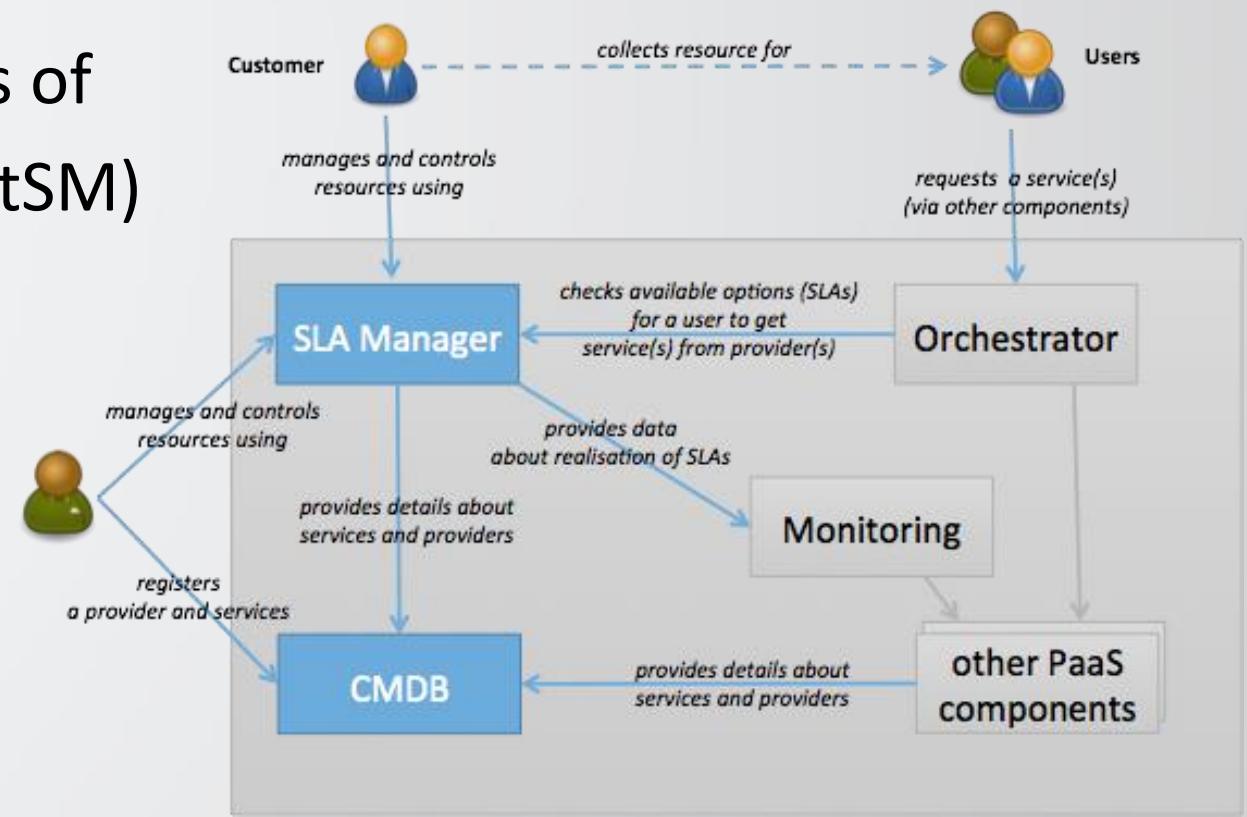


- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- **CMDB**
- Managed Services/Application (MSA) Deployment Service
- OneData

# CMDB Service



- Indigo Configuration Management DB (CMDB)
  - Need to know what are INDIGO providers and services
  - need to register specific data
- Idea comes from good practices of IT Service Management (ITIL, FitSM)



# Index

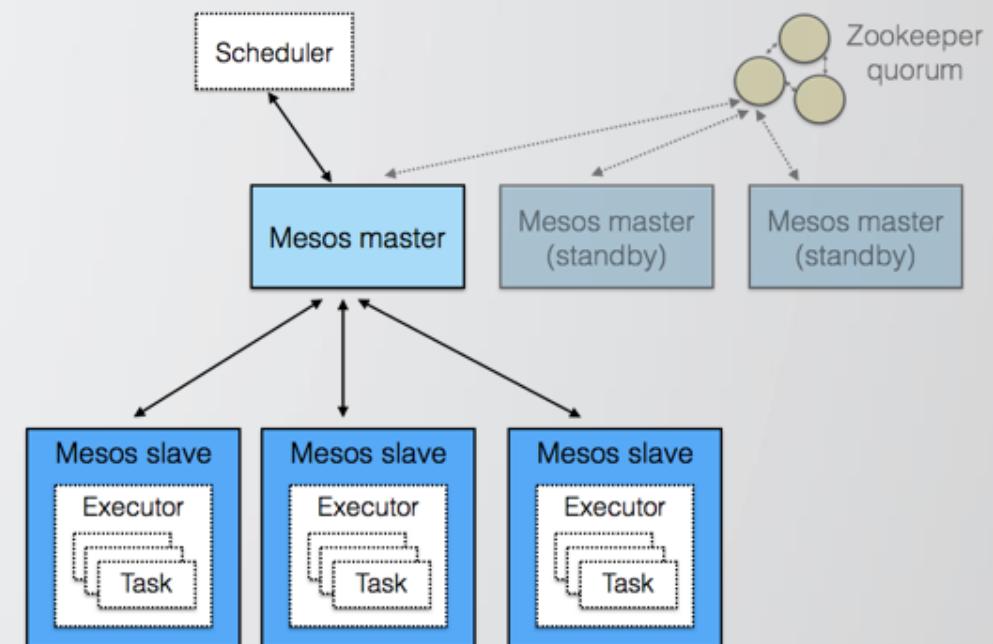


- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- **Managed Services/Application (MSA) Deployment Service**
- OneData

# Managed Services/Application (MSA) Deployment Service



- This service is in charge of scheduling, spawning, executing and monitoring applications and services on a distributed infrastructure.
- The core of this component consists of an elastic Apache Mesos cluster with slave nodes dynamically provisioned and distributed on the IaaS sites.
- Apache Mesos provides efficient resource isolation and sharing across distributed applications (frameworks).



# Index



- Indigo PaaS Overview
- TOSCA
- IM
- CLUES
- Orchestrator
- Cloud Provider Ranker
- Monitoring Service
- SLAM
- Accounting Service
- CMDB
- Managed Services/Application (MSA) Deployment Service
- **OneData**

# Data in multi-cloud environments



# Onedata – Features (I)



- Unified vision of geographically distributed data set
- Data affinity
  - Computation jobs started on resources close to data.
- Federated data access
  - Interoperability and OpenData
- Optimization and Data on the fly
  - when data is not staged

# Onedata – Features (II)



- Multi-protocol transparent access to data in multi-cloud environments
  - Transparent access your data and create new one in multi-cloud environments
  - Care less about data locality, all your data are accessible wherever you go
  - Use many protocols to access the same data
- Heterogeneity of storage technologies
  - Use the same data access protocols (up to your choice) wherever you go
  - Pass-through problems of selection right storage technology to data centres operators
  - Avoid cloud vendor locking

# Onedata – Features (III)



- Replica management
  - Replicate files on demand and on the fly without any additional effort
  - Migrate data between sites on demand with simple API interface
  - Easily check location of your data through GUI or API
- Easy data sharing without borders
  - Share large scale data collection with other communities
  - Enable your data to be shared in cross-federation scenarios
  - Bring your data and tools as building blocks to European Open Science Cloud

# Data sharing without borders



- Team-sharing
  - For groups
  - For individuals
  - Using tokens
- Cross-community data sharing
- Instant and ad-hoc data sharing
- *Thanks to effort supported by EGI Engage:*
  - Open Data Publication
  - Handles (DOI) support
  - OAI-PMH

A screenshot of the ONE DATA interface. On the left, there's a sidebar with icons for Data, Spaces (highlighted in red), Groups, Tokens, and Providers. The main area has a header "ONE DATA" with "SPACES" and "Create" buttons. Below is a list of data spaces: "Astronomy Da...", "Users", "Groups", "Big Data Exper...", and "Cancer Data". To the right, there's a "USERS" table with columns for "Invite user", "VIEW SPACE", "MODIFY SPACE", "SET PRIVILEGES", and "REMOVE SPACE". The table lists users: adam, iza, ola, and orzech, each with corresponding checkboxes for each privilege level.

# Onedata – Features (IV)

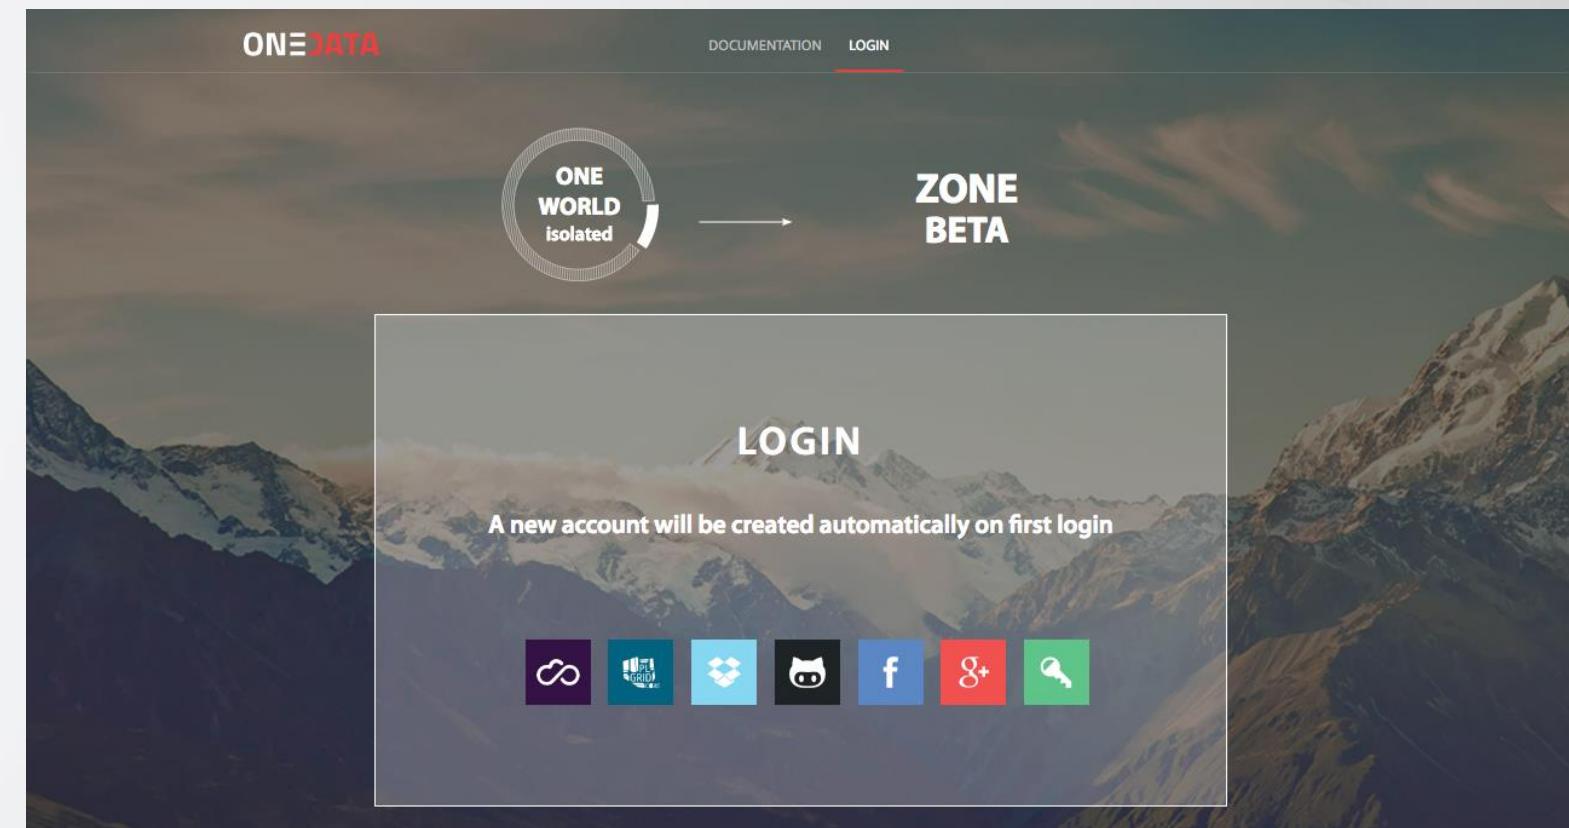


- Metadata management integrated with data management platform
  - Work with data and metadata in one system – avoiding problems of consistency
  - Monitor metadata data changes through API in order to feed external custom systems
- Flexible authentication and authorization
  - Control who knows about your data
  - Control who can access data on a single file level
- Easy integration using API with external tools
  - Integrate external tools using rich API interfaces with data management platform building more complex environment for data processing

# Flexible authentication and authorization



- Integrated with Indigo IAM
- Pluggable methods of authentication per zone
- Multi level of access control
- ACL on files and directories
- Group management
- Token based authentication (macaroons)
- X.509 in prep.



# QUESTIONS?



Please visit:

**[www.indigo-datacloud.eu](http://www.indigo-datacloud.eu)**